

Answers to questions in

Lab 2: Edge detection & Hough transform

Name: Alex Gunnarsson Program: TCSCM

Instructions: Complete the lab according to the instructions in the notes and respond to the questions stated below. Keep the answers short and focus on what is essential. Illustrate with figures only when explicitly requested.

Good luck!

Question 1: What do you expect the results to look like and why? Compare the size of *dxttools* with the size of *tools*. Why are these sizes different?

Answers: Since the Sobel operator approximates the first order derivative in the x and y directions, we get highlights of where the image changes rapidly in either orthogonal direction.



The size of the original image is 256x256 while both dxttools and dytools are 254x254. This is because we cannot calculate the values along the edges so we lose 1 pixel on all sides since the matrix is 3x3. The “valid” parameter ensures this and does not add padding to help in calculating the borders.

Question 2: Is it easy to find a threshold that results in thin edges? Explain why or why not!

Answers: Not really, there is a bit of noise in both images that is either detected as an edge, or the real edge is not fully detected. This is also true for the higher complexity image.


Raw




gradmagntools




Threshold = 50




Threshold = 100




Threshold = 150

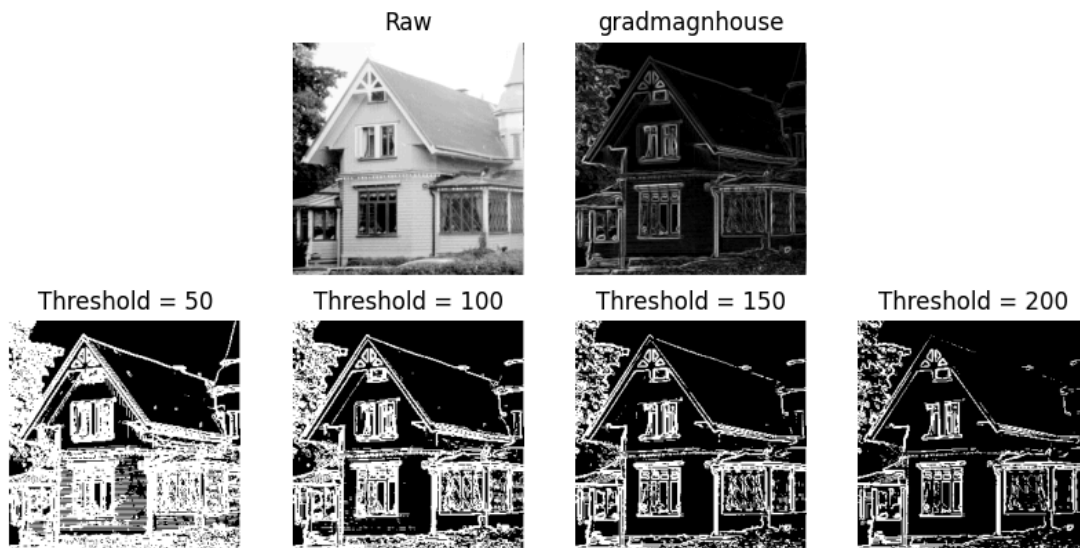


Threshold = 200




Histogram gradmagntools



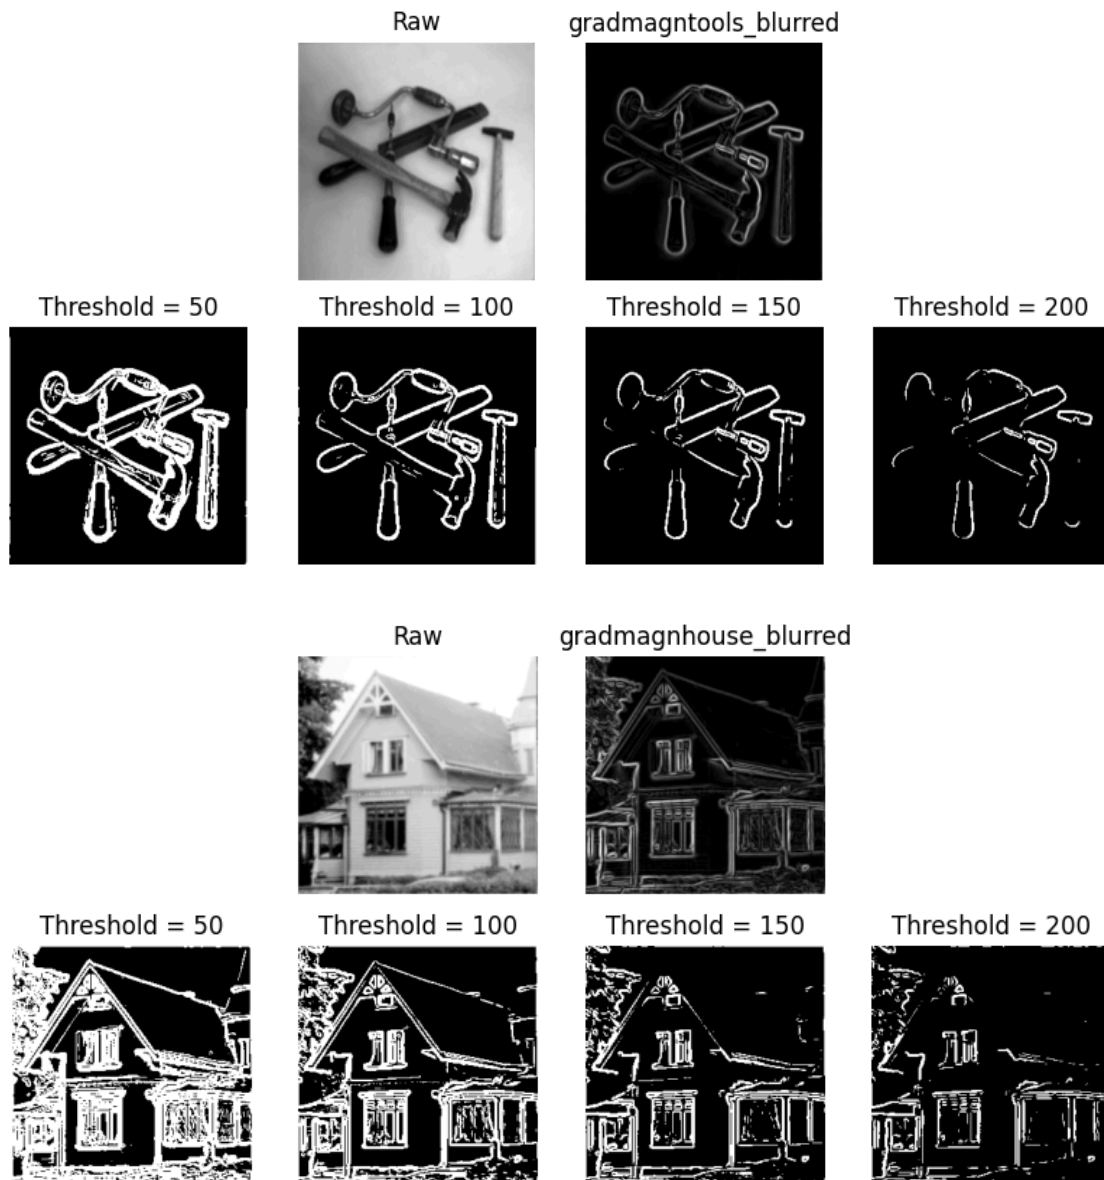


Histogram gradmagnhouse



Question 3: Does smoothing the image help to find edges?

Answers: Perhaps slightly, although it mostly seems to shift the optimal threshold value, while the optimal threshold image remains pretty much the same. The images below were blurred with a Gaussian kernel with $\sigma=1$.




Question 4: What can you observe? Provide explanation based on the generated images.

Answers: The goal is to find where an edge is by identifying where there are drastic changes in the image, L . Quite naturally, this should be the one point where the image is changing the most, so it should be where L_v (the first derivative) is the largest, which is given by $L_{vv}=0$, and to make sure that L_v is the largest, not the smallest, which is given by $L_{vvv}<0$.



Question 5: Assemble the results of the experiment above into an illustrative collage with the *subplot* command. Which are your observations and conclusions?

Answers: We can see that we get where more general regions of where relevant edges may be.




Question 6: How can you use the response from L_{vv} to detect edges, and how can you improve the result by using L_{vvv} ?

Answers: By combining the two, such that an edge is where $L_{vv}=0$ and $L_{vvv}<0$.

Question 7: Present your best results obtained with *extractedge* for *house* and *tools*.


Answers:

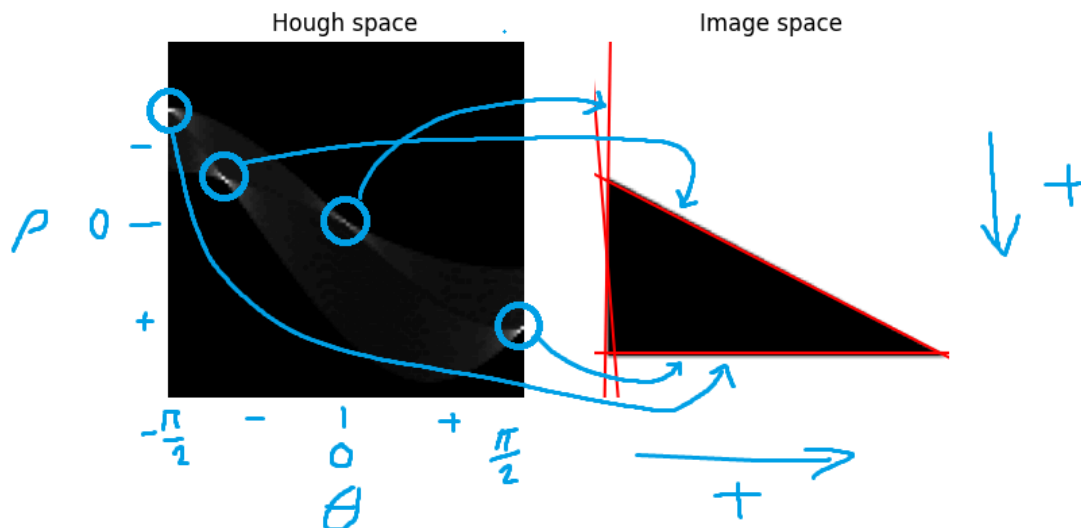


Question 8: Identify the correspondences between the strongest peaks in the accumulator and line segments in the output image. Doing so convince yourself that the implementation is correct. Summarize the results of in one or more figures.

Answers: The parametrization yields

$y = \frac{\rho - x \cdot \cos(\theta)}{\sin(\theta)} = \frac{\rho}{\sin(\theta)} - x \cdot \frac{\cos(\theta)}{\sin(\theta)}$ so the line will have intersection at $x = 0$ and $y = \frac{\rho}{\sin(\theta)}$ with a slope of $-\frac{\cos(\theta)}{\sin(\theta)}$, which will look like:






Question 9: How do the results and computational time depend on the number of cells in the accumulator?

Answers: The computational time increases, but is more dependent on the number of theta values rather than the number of rho values, since thetas are iterated through, while rhos are just searched for and found.

Below are the results from the given images.

With `houghedgeline(img1, 8, 35, 300, 300, 20, 2):`



With `houghedgeline(img2, 2, 45, 500, 500, 20, 2):`

Edgecurves



Hough space





Image space



With `houghedgeline(img3, 5, 35, 500, 500, 20, 2):`

Edgecurves



Hough space




Image space



And some variations:


`houghedgeline(img2, 2, 30, 800, 100, 20, 2):`

Image space



`houghedgeline(img3, 1, 70, 1000, 200, 40, 2):`

Image space



Question 10: How do you propose to do this? Try out a function that you would suggest and see if it improves the results. Does it?

Answers: We can weight them according to the magnitude of the gradient by prioritising larger magnitudes and counting their vote as more valuable. Magnitudes closer to the threshold should be counted less, while magnitudes farther away from the threshold should be counted equally much. Therefore, the logarithm of the difference between the magnitude and the threshold makes sense, but as we see below, it didn't make much of an actual difference. The only visible difference is that it found an extra edge on the top of the hammer in "tools".

Edgecurves



Hough space





Image space




Edgecurves



Hough space



Image space



Edgecurves



Hough space




Image space



Image space




Image space

